

Assignment 3 - Due Tuesday, November 23, 2021

The skeleton for this assignment is available at the master branch of <https://git.uwaterloo.ca/ece650-f2021/skeleton> in directory `a3`. Follow the instructions in Assignment 0 to correctly fetch and merge the files from the skeleton!

For this assignment, you are to:

- Modify the output format of your Python script from Assignment 1 to match the input format of your C++ program from Assignment 2.
- Modify your C++ program from Assignment 2 to output the input graph on standard output.
- Write a program in C++ to generate random input for your Python script.
- Write a driver program in C++ that uses Inter-Process Communication (IPC) to link the output of the random input generator to the input of the Python script, and the output of the Python script to the input of the C++ program from Assignment 2.

Resources

The online book, “Operating Systems: Three Easy Pieces”, will be useful for this assignment and beyond. For this assignment, you might find Chapter 4 on Processes and Chapter 5 on Process API the most useful. The book is available online at <http://pages.cs.wisc.edu/~remzi/OSTEP/>.

You might also want to consult “Advanced Linux Programming”, especially Chapter 3 on Processes and Chapter 5.4 on Pipes. The book is available from <https://github.com/MentorEmbedded/advancedlinuxprogramming/blob/gh-pages/alp-folder/advanced-linux-programming.pdf>

Sample Run

You should name the driver’s executable `ece650-a3`. In the following, we assume that “\$” is the shell command-prompt.

```
$ cd build
$ cmake ../
$ make install
$ cd ./run/bin
$ ./ece650-a3 -s 5 -n 4 -l 5
V 8
E {<1,3>,<1,4>,<1,5>,<2,4>,<5,8>,<6,3>,<6,7>}
s 3 5
3-1-5
```

In the above, the first three lines make your executable, and run the driver program with some command-line arguments. Then, the lines “V = ...”, “E = ...”, and “3-1-5” are output. The input the user provided to your program via stdin is “s 3 5”.

Input, Output, and Error

Your program should take input from stdin, and output to stdout. Errors should be output to stderr. Errors should always start with “Error:” followed by a brief description. All your processes should terminate gracefully (and quietly) once you see EOF at stdin. Your program should not generate any extraneous output; for example, do not print out prompt strings such as “please enter input” and things like that.

As the example above indicates, there are two kinds of input the user provides. One is via the command-line arguments, with switches such as `-s` and `-n`. This is done only once, when your program is first started. The other type of input is the ‘s’ command on stdin, which may be issued

repeatedly, just as in Assignment 2. For the 's' command, your program should output a shortest path.

We will not test your program for format errors in the input. That is, the command-line arguments, whenever specified, will be formatted correctly, and the s input will also be formatted correctly. Of course, we may omit command-line arguments (see below for what to do in such cases), and specify vertex IDs with s that do not exist, or between whom a path does not exist. The latter two cases should cause your program to report an error.

Marking

- Does not compile/make/crashes: automatic 0
- Your program runs, awaits input and does not crash on input: + 20
- Run in default mode: + 15
- Error check arguments: + 5
- Test functionality: + 20
- Test rgen: + 20
- Replace rgen: + 20

CMake

As discussed below under “Submission Instructions”, you should use a CMakeLists.txt file to build your project. We will build your project using the following sequence:

```
cd a3 && mkdir build && cd build && cmake ../ && make install
```

If your code is not compiled from scratch (i.e., from the C++ sources), you get an automatic 0.

Note that we are using `make install` instead of `make`. The `install` target instructs `make` to copy all of the binaries (both Python and C++) into directory `./run/bin`. This ensures that all the executable files (including the Python program) are in the same directory.

Submission Instructions

You should place all your files in the `a3` directory in your GitLab repository. The directory should contain:

- All your C++ and Python source-code files. IMPORTANT NOTE: the executable for your random generator must be named `rgen`. The reason is that part of our testing will replace your generator with ours.
- A `CMakeLists.txt`, that builds all of the C++ executables: `rgen`, `ece650-a2`, and `ece650-a3`.
- A file `user.yml` that includes your name, WatIAM, and student number.

See `README.md` for any additional information.

The submitted files should be in the `a3` directory of the `master` branch of your repository.

File names

There are two file names that are important. One is your main executable. This must be named `ece650-a3`. Our marking script will simply try and run this after building your project. The other executable file whose name is important is your random input generator. The executable for this must be named `rgen`. The reason is that for some of our tests, we will replace your `rgen` with ours.

Things to be done

Python script

You should edit your Python script from Assignment 1 so that its output which is the specification of the graph is formatted to match the input format of your C++ program from Assignment 2.

Think of it as though your Python script provides command-line input to the C++ program. The way to do this is to simply map each vertex to an index in $[1, n]$ (where n is the number of vertices), and rename your edges accordingly.

Also, the only output that your Python script should produce to stdout is in response to “gg”, for which it outputs the specification of the graph (i.e., V and E). Error output should go to stderr.

Cpp program from Assignment 2

You should edit it to print the graph that it has read before accepting an `s` command. This is necessary so that you know what graph has been produced and which vertices are available.

Random input generator

Your random input generator `rngen` should generate random street specifications as input for your Python script. It takes four command-line arguments. All are optional.

- `-s k` — where k is an integer ≥ 2 . The number of streets should be a random integer in $[2, k]$. If this option is not specified, you should use a default of $k = 10$; that is, the number of streets should be a random integer in $[2, 10]$.
- `-n k` — where k is an integer ≥ 1 . The number of line-segments in each street should be a random integer in $[1, k]$. Default: $k = 5$.
- `-l k` — where k is an integer ≥ 5 . Your process should wait a random number w seconds, where w is in $[5, k]$ before generating the next (random) input. Default: $k = 5$.
- `-c k` — where k is an integer ≥ 1 . Your process should generate (x, y) coordinates such that every x and y value is in the range $[-k, k]$. For example, if $k = 15$, all of your coordinate values should be integers between -15 and 15 . Default: $k = 20$.

Your program should generate a specification of streets in the format that your Python script expects (see Assignment 1). You can name the streets whatever you want. You should ensure that your input does not have errors. For example, if you generate a line-segment that overlaps with a line-segment (across all streets) generated earlier, you should regenerate that line-segment. Similarly, you should not have any zero-length line segments.

Also, note that your random generator could go into an infinite loop looking for a valid specification. You should disallow this from happening by limiting the number of tries. That is, if your random generator fails to generate a valid specification for a continuous A number of attempts, it should `exit()` with an error message reported on stderr. A reasonable A to adopt may be 25. Whatever A you adopt, your error message should mention it. That is, your error message should be something like, “Error: failed to generate valid input for 25 simultaneous attempts”.

Before adding a new street graph specification to your Python script, your generator should issue “rm” commands to your Python script to remove all previous streets and replace them with the new street specification.

After generating the input, the generator must issue the “gg” command.

Thus, a typical interaction of the random generator is as follows:

1. issue enough `rm` commands to clear any existing street database;
2. issue `add` commands to add new streets satisfying the specification;
3. issue a `gg` command;
4. wait for specified number of seconds and repeat.

You must use `/dev/urandom` as the source of your random data. Also, see under “Submission Instructions” how your executable for the random generator must be named.

Driver

Your driver program has the overall control. You have at least three programs that run concurrently: (1) the random generator, (2) your Python script from Assignment 1, and, (3) your program from Assignment 2 that takes a graph-specification as input and computes shortest paths. Your driver program should `fork()` two processes and `exec()` two of those programs, and then `exec()` the third (so it turns into a process that corresponds to the third program). It should set up all IPC appropriately beforehand.

It should send normal output to `stdout`, error output to `stderr`, and take input from `stdin`. As we mention above, the only input it takes are “s” commands, that ask for a shortest path between vertices.

Gotcha warning: note that I say “you have at least three programs that run concurrently.” You may need more than three that run concurrently. (Why?)

Sample code

The skeleton repository for the assignment contains some sample code. You should replace all sample code with your own code. Including replacing the Python code with your solution from Assignment 1, and C++ code from your solution to Assignment 2.